

Advice for applying machine learning

1. Deciding what to try next

Debugging a learning algorithm

You've implemented regularized linear regression on housing prices

$$\rightarrow J(\vec{w}, b) = \frac{1}{2m} \sum_{i=1}^m (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n w_j^2$$

But it makes unacceptably large errors in predictions. What do you try next?

- \rightarrow Get more training examples
- \rightarrow Try smaller sets of features
- \rightarrow Try getting additional features
- \rightarrow Try adding polynomial features ($x_1^2, x_2^2, x_1x_2, etc$)
- \rightarrow Try decreasing λ
- \rightarrow Try increasing λ

Machine learning diagnostic


Diagnostic:

A test that you run to gain insight into what is/isn't working with a learning algorithm, to gain guidance into improving its performance.

Diagnostics can take time to implement but doing so can be a very good use of your time.

2. Evaluating a model

Evaluating your model



Evaluating your model

Dataset:

	size	price	
70%	2104	400	$\xrightarrow{\text{+training set}}$ $m_{\text{train}} = \text{no. training examples}$ $= 7$
	1600	330	
	2400	369	
	1416	232	
	3000	540	
	1985	300	
	1534	315	
30%	1427	199	$\xrightarrow{\text{+test set}}$ $m_{\text{test}} = \text{no. test examples}$ $= 3$
	1380	212	
	1494	243	

$(x^{(1)}, y^{(1)})$
 $(x^{(2)}, y^{(2)})$
 \vdots
 $(x^{(m_{\text{train}})}, y^{(m_{\text{train}})})$

$(x_{\text{test}}^{(1)}, y_{\text{test}}^{(1)})$
 \vdots
 $(x_{\text{test}}^{(m_{\text{test}})}, y_{\text{test}}^{(m_{\text{test}})})$

Train/test procedure for linear regression (with squared error cost)

Fit parameters by minimizing cost function $J(\vec{w}, b)$

$$\rightarrow J(\vec{w}, b) = \left[\frac{1}{2m_{\text{train}}} \sum_{i=1}^{m_{\text{train}}} (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m_{\text{train}}} \sum_{j=1}^n w_j^2 \right]$$

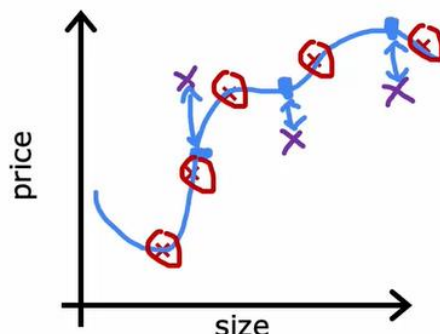
Compute test error:

$$J_{\text{test}}(\vec{w}, b) = \frac{1}{2m_{\text{test}}} \left[\sum_{i=1}^{m_{\text{test}}} (f_{\vec{w}, b}(\vec{x}_{\text{test}}^{(i)}) - y_{\text{test}}^{(i)})^2 \right] \quad \text{with } \cancel{\sum_{j=1}^n w_j^2}$$

Compute training error:

$$J_{\text{train}}(\vec{w}, b) = \frac{1}{2m_{\text{train}}} \left[\sum_{i=1}^{m_{\text{train}}} (f_{\vec{w}, b}(\vec{x}_{\text{train}}^{(i)}) - y_{\text{train}}^{(i)})^2 \right]$$

Train/test procedure for linear regression (with squared error cost)



$\times = \text{+train}$

$\times = \text{+test}$

$J_{\text{train}}(\vec{w}, b)$ will be low

$J_{\text{test}}(\vec{w}, b)$ will be high

Train/test procedure for classification problem

0/1

Fit parameters by minimizing $J(\vec{w}, b)$ to find \vec{w}, b

E.g.,

$$J(\vec{w}, b) = -\frac{1}{m_{train}} \sum_{i=1}^{m_{train}} \left[y^{(i)} \log(f_{\vec{w}, b}(\vec{x}^{(i)})) + (1 - y^{(i)}) \log(1 - f_{\vec{w}, b}(\vec{x}^{(i)})) \right] + \frac{\lambda}{2m_{train}} \sum_{j=1}^n w_j^2$$

Compute test error:

$$J_{test}(\vec{w}, b) = -\frac{1}{m_{test}} \sum_{i=1}^{m_{test}} \left[y_{test}^{(i)} \log(f_{\vec{w}, b}(\vec{x}_{test}^{(i)})) + (1 - y_{test}^{(i)}) \log(1 - f_{\vec{w}, b}(\vec{x}_{test}^{(i)})) \right]$$

Compute train error:

$$J_{train}(\vec{w}, b) = -\frac{1}{m_{train}} \sum_{i=1}^{m_{train}} \left[y_{train}^{(i)} \log(f_{\vec{w}, b}(\vec{x}_{train}^{(i)})) + (1 - y_{train}^{(i)}) \log(1 - f_{\vec{w}, b}(\vec{x}_{train}^{(i)})) \right]$$

Train/test procedure for classification problem

fraction of the test set and the fraction of the train set that the algorithm has misclassified.

$$\hat{y} = \begin{cases} 1 & \text{if } f_{\vec{w}, b}(\vec{x}^{(i)}) \geq 0.5 \\ 0 & \text{if } f_{\vec{w}, b}(\vec{x}^{(i)}) < 0.5 \end{cases}$$


count $\hat{y} \neq y$

$J_{test}(\vec{w}, b)$ is the fraction of the test set that has been misclassified.

$J_{train}(\vec{w}, b)$ is the fraction of the train set that has been misclassified.

3. Model selection and training/cross validation/test sets

Model selection (choosing a model)



$x = x_1$

$$f_{\vec{w}, b}(\vec{x}) = w_1 x + w_2 x^2 + w_3 x^3 + w_4 x^4 + b$$

Once parameters \vec{w}, b are fit to the training set, the training error $J_{train}(\vec{w}, b)$ is likely lower than the actual generalization error.

$J_{test}(\vec{w}, b)$ is better estimate of how well the model will generalize to new data compared to $J_{train}(\vec{w}, b)$.

Model selection (choosing a model)

- $d=1$ 1. $f_{\vec{w},b}(\vec{x}) = w_1x + b \rightarrow w^{<1>}, b^{<1>} \rightarrow J_{test}(w^{<1>}, b^{<1>})$
 $d=2$ 2. $f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + b \rightarrow w^{<2>}, b^{<2>} \rightarrow J_{test}(w^{<2>}, b^{<2>})$
 $d=3$ 3. $f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + w_3x^3 + b \rightarrow w^{<3>}, b^{<3>} \rightarrow J_{test}(w^{<3>}, b^{<3>})$
 \vdots
 $d=10$ 10. $f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + \dots + w_{10}x^{10} + b \rightarrow J_{test}(w^{<10>}, b^{<10>})$
- Choose $w_1x + \dots + w_5x^5 + b$ $d=5$ $J_{test}(w^{<5>}, b^{<5>})$

How well does the model perform? Report test set error $J_{test}(w^{<5>}, b^{<5>})$?
 The problem: $J_{test}(w^{<5>}, b^{<5>})$ is likely to be an optimistic estimate of generalization error (ie. $J_{test}(w^{<5>}, b^{<5>}) < \text{generalization error}$). Because an extra parameter d (degree of polynomial) was chosen using the test set.

w, b are overly optimistic estimate of generalization error on training data.

Training/cross validation/test set

size	price		validation set	development set
2104	400	} \rightarrow training set 60%	$(x^{(1)}, y^{(1)})$ \vdots $(x^{(m_{train})}, y^{(m_{train})})$	\rightarrow dev set
1600	330			
2400	369			
1416	232			
3000	540			
1985	300			
1534	315	} cross validation 20%	$(x_{cv}^{(1)}, y_{cv}^{(1)})$ \vdots $(x_{cv}^{(m_{cv})}, y_{cv}^{(m_{cv})})$	\rightarrow
1427	199			
1380	212			
1494	243	} test set 20%	$(x_{test}^{(1)}, y_{test}^{(1)})$ \vdots $(x_{test}^{(m_{test})}, y_{test}^{(m_{test})})$	\rightarrow

$m_{train} = 6$
 $m_{cv} = 2$
 $m_{test} = 2$

Training/cross validation/test set

Training error:
$$J_{train}(\vec{w}, b) = \frac{1}{2m_{train}} \left[\sum_{i=1}^{m_{train}} (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)})^2 \right]$$

Cross validation error:
$$J_{cv}(\vec{w}, b) = \frac{1}{2m_{cv}} \left[\sum_{i=1}^{m_{cv}} (f_{\vec{w},b}(\vec{x}_{cv}^{(i)}) - y_{cv}^{(i)})^2 \right] \quad (\text{validation error, dev error})$$

Test error:
$$J_{test}(\vec{w}, b) = \frac{1}{2m_{test}} \left[\sum_{i=1}^{m_{test}} (f_{\vec{w},b}(\vec{x}_{test}^{(i)}) - y_{test}^{(i)})^2 \right]$$


Model selection

$$\begin{array}{ll}
 d=1 & 1. f_{\vec{w},b}(\vec{x}) = w_1x + b \quad w^{(1)}, b^{(1)} \rightarrow J_{cv}(w^{(1)}, b^{(1)}) \\
 d=2 & 2. f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + b \quad \rightarrow J_{cv}(w^{(2)}, b^{(2)}) \\
 d=3 & 3. f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + w_3x^3 + b \\
 \vdots & \vdots \\
 d=10 & 10. f_{\vec{w},b}(\vec{x}) = w_1x + w_2x^2 + \dots + w_{10}x^{10} + b \quad J_{cv}(w^{(10)}, b^{(10)})
 \end{array}$$

→ Pick $w_1x + \dots + w_4x^4 + b$ ($J_{cv}(w^{(4)}, b^{(4)})$)

Estimate generalization error using test the set: $J_{test}(w^{(4)}, b^{(4)})$

Model selection – choosing a neural network architecture



Pick $w^{(2)}, b^{(2)}$

Estimate generalization error using the test set: $J_{test}(w^{(2)}, b^{(2)})$

4. Practice quiz

1.

In the context of machine learning, what is a diagnostic?

- ☐ This refers to the process of measuring how well a learning algorithm does on a test set (data that the algorithm was not trained on).
- ☐ An application of machine learning to medical applications, with the goal of diagnosing patients' conditions.
- ☒ A test that you run to gain insight into what is/isn't working with a learning algorithm.
- ☐ A process by which we quickly try as many different ways to improve an algorithm as possible, so as to see what works.

✓ Correct

Yes! A diagnostic is a test that you run to gain insight into what is/isn't working with a learning algorithm, to gain guidance into improving its performance.

2.

True/False? It is always true that the better an algorithm does on the training set, the better it will do on generalizing to new data.

- ☒ False
- ☐ True

✓ **Correct**

Actually, if a model overfits the training set, it may not generalize well to new data.

For a classification task; suppose you train three different models using three different neural network architectures. Which data do you use to evaluate the three models in order to choose the best one?

- ☐ All the data -- training, cross validation and test sets put together.
- ☐ The training set
- ☐ The test set
- ☒ The cross validation set

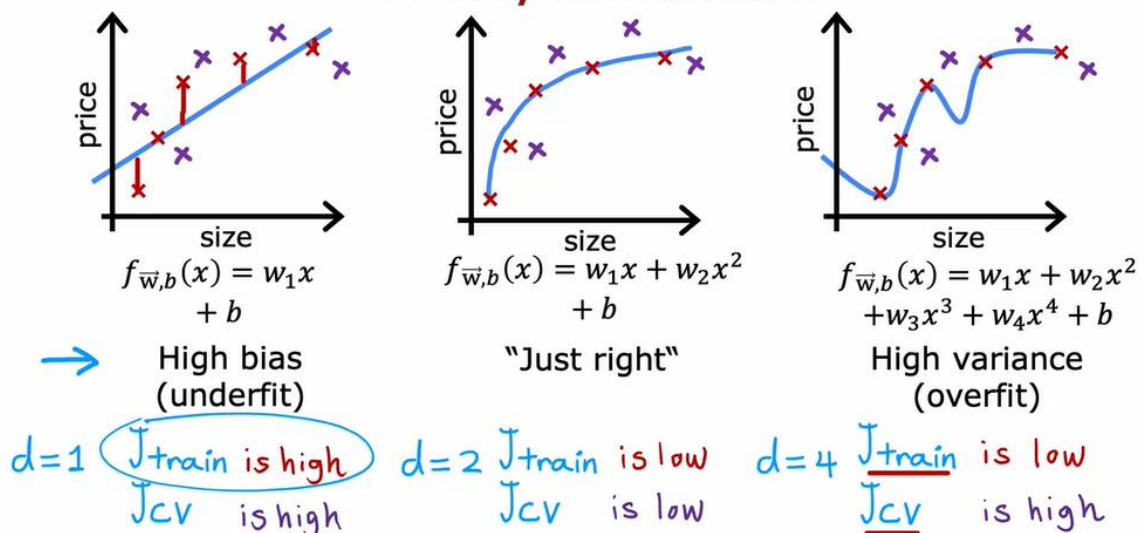
✓ **Correct**

Correct. Use the cross validation set to calculate the cross validation error on all three models in order to compare which of the three models is best.


Bias and variance

1. Diagnosing bias and variance

Bias/variance




Understanding bias and variance



Diagnosing bias and variance

How do you tell if your algorithm has a bias or variance problem?



High bias (underfit)

→ J_{train} will be high
($J_{train} \approx J_{cv}$)

High variance (overfit)

→ $J_{cv} \gg J_{train}$
(J_{train} may be low)

High bias and high variance


→ J_{train} will be high
and $J_{cv} \gg J_{train}$

2. Regularization and bias/variance

Linear regression with regularization

Model: $f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + b$

$$J(\vec{w}, b) = \frac{1}{2m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n w_j^2$$



Choosing the regularization parameter λ


Model: $f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + b$

- \rightarrow 1. Try $\lambda = 0 \rightarrow \min_{\vec{w},b} J(\vec{w},b) \rightarrow w^{<1>}, b^{<1>} \rightarrow J_{cv}(w^{<1>}, b^{<1>})$
 \rightarrow 2. Try $\lambda = 0.01 \rightarrow w^{<2>}, b^{<2>} \rightarrow J_{cv}(w^{<2>}, b^{<2>})$
 \rightarrow 3. Try $\lambda = 0.02 \rightarrow w^{<3>}, b^{<3>} \rightarrow J_{cv}(w^{<3>}, b^{<3>})$
 \rightarrow 4. Try $\lambda = 0.04 \rightarrow w^{<5>}, b^{<5>} \rightarrow J_{cv}(w^{<5>}, b^{<5>})$
 \rightarrow 5. Try $\lambda = 0.08$
 \vdots
 \rightarrow 12. Try $\lambda \approx 10 \rightarrow w^{<12>}, b^{<12>} \rightarrow J_{cv}(w^{<12>}, b^{<12>})$

Pick $w^{<5>}, b^{<5>}$

Report test error: $J_{test}(w^{<5>}, b^{<5>})$

Bias and variance as a function of regularization parameter λ



3. Establishing a baseline level of performance

Speech recognition example



Human level performance : 10.6%
 Training error J_{train} : 10.8%
 Cross validation error J_{cv} : 14.8%

Differences: 0.2% (between Human and Training), 4.0% (between Training and Cross validation)



Establishing a baseline level of performance

What is the level of error you can reasonably hope to get to?

- • Human level performance
- • Competing algorithms performance
- • Guess based on experience

Bias/variance examples



Baseline performance	: 10.6%	↓ 0.2%	10.6%	↓ 4.4%	10.6%	↓ 4.4%
Training error (J_{train})	: 10.8%	↓ 4.0%	15.0%	↓ 0.5%	15.0%	↓ 4.7%
Cross validation error (J_{cv})	: 14.8%		15.5%		19.7%	
		high variance	high bias	high bias	high variance	

4. Learning curves

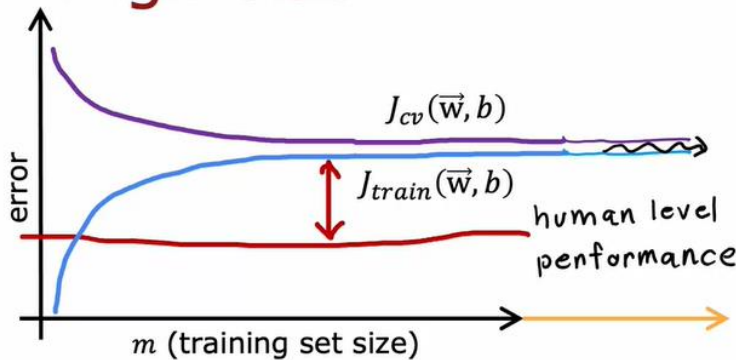
Learning curves

J_{train} = training error


J_{cv} = cross validation error



High bias




$$f_{\vec{w},b}(x) = w_1x + b$$




if a learning algorithm suffers from high bias, getting more training data will not (by itself) help much.

High variance



$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + b$$

(with small λ)



If a learning algorithm suffers from high variance, getting more training data is likely to help.

5. Deciding what next revisited

Debugging a learning algorithm

You've implemented regularized linear regression on housing prices

$$J(\vec{w}, b) = \frac{1}{2m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n w_j^2$$

But it makes unacceptably large errors in predictions. What do you try next?

- Get more training examples
- Try smaller sets of features x, x^2, x^3, x^4, \dots
- Try getting additional features
- Try adding polynomial features $(x_1^2, x_2^2, x_1x_2, \text{etc})$
- Try decreasing λ
- Try increasing λ

fixes high variance
 fixes high variance
 fixes high bias
 fixes high bias
 fixes high bias
 fixes high variance

The bias variance tradeoff

$$f_{\vec{w},b}(x) = w_1x + b$$

Simple model


High bias

$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + b$$

$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + b$$

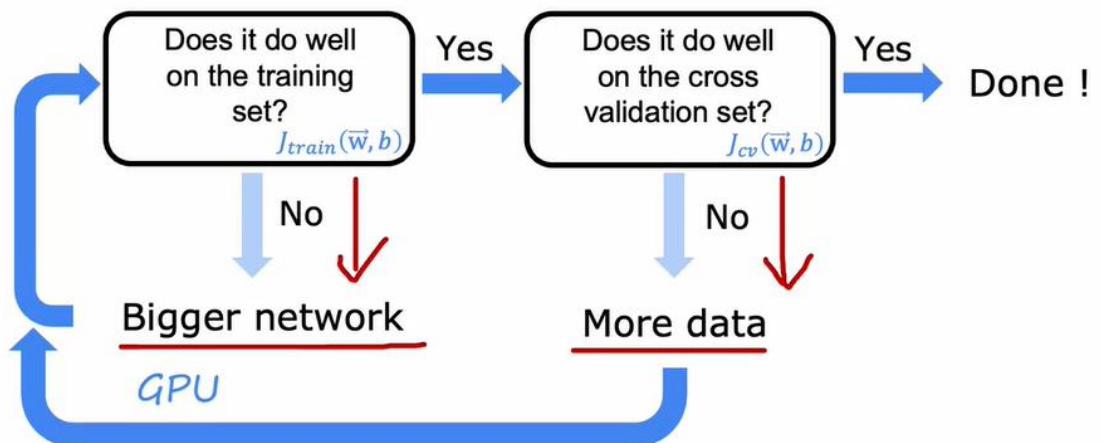
Complex model

High variance




Neural networks and bias variance

Large neural networks are low bias machines



Neural networks and regularization



A large neural network will usually do as well or better than a smaller one so long as regularization is chosen appropriately.

Neural network regularization

$$J(\mathbf{W}, \mathbf{B}) = \underbrace{\frac{1}{m} \sum_{i=1}^m L(f(\vec{x}^{(i)}), y^{(i)})}_{\text{training error}} + \underbrace{\frac{\lambda}{2m} \sum_{\text{all weights } \mathbf{W}} (w^2)}_{\text{regularization term}} \quad b$$

Unregularized MNIST model

```
layer_1 = Dense(units=25, activation="relu")
layer_2 = Dense(units=15, activation="relu")
layer_3 = Dense(units=1, activation="sigmoid")
model = Sequential([layer_1, layer_2, layer_3])
```

Regularized MNIST model

```
layer_1 = Dense(units=25, activation="relu", kernel_regularizer=L2(0.01))
layer_2 = Dense(units=15, activation="relu", kernel_regularizer=L2(0.01))
layer_3 = Dense(units=1, activation="sigmoid", kernel_regularizer=L2(0.01))
model = Sequential([layer_1, layer_2, layer_3])
```

6. Practice quiz

If the model's cross validation error J_{cv} is much higher than the training error J_{train} , this is an indication that the model has...

- ☒ high variance
- ☐ Low variance
- ☐ high bias
- ☐ Low bias

✓ Correct

When $J_{cv} \gg J_{train}$ (whether J_{train} is also high or not, this is a sign that the model is overfitting to the training data and performing much worse on new examples.

Which of these is the best way to determine whether your model has high bias (has underfit the training data)?

- ☐ See if the cross validation error is high compared to the baseline level of performance
- ☒ Compare the training error to the baseline level of performance
- ☐ See if the training error is high (above 15% or so)
- ☐ Compare the training error to the cross validation error.

✓ Correct

Correct. If comparing your model's training error to a baseline level of performance (such as human level performance, or performance of other well-established models), if your model's training error is much higher, then this is a sign that the model has high bias (has underfit).

You find that your algorithm has high bias. Which of these seem like good options for improving the algorithm's performance? Hint: two of these are correct.

- ☒ Decrease the regularization parameter λ (lambda)

☒ **Correct**

Correct. Decreasing regularization can help the model better fit the training data.

- ☒ Collect additional features or add polynomial features

☒ **Correct**

Correct. More features could potentially help the model better fit the training examples.

- ☐ Remove examples from the training set

- ☐ Collect more training examples

You find that your algorithm has a training error of 2%, and a cross validation error of 20% (much higher than the training error). Based on the conclusion you would draw about whether the algorithm has a high bias or high variance problem, which of these seem like good options for improving the algorithm's performance? Hint: two of these are correct.

- ☒ Increase the regularization parameter λ

☒ **Correct**

Yes, the model appears to have high variance (overfit), and increasing regularization would help reduce high variance.

- ☒ Collect more training data

☒ **Correct**

Yes, the model appears to have high variance (overfit), and collecting more training examples would help reduce high variance.


- ☐ Decrease the regularization parameter λ

- ☐ Reduce the training set size

Machine learning development process

1. Iterative loop of ML development

Iterative loop of ML development



Building a spam classifier

Supervised learning: \vec{x} = features of email
 y = spam (1) or not spam (0)

Features: list the top 10,000 words to compute $x_1, x_2, \dots, x_{10,000}$

$\vec{x} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ \vdots \end{bmatrix}$	$\begin{bmatrix} a \\ andrew \\ buy \\ deal \\ discount \\ \vdots \end{bmatrix}$	<div>From: cheapsales@buystufffromme.com To: <u>Andrew</u> Ng Subject: <u>Buy</u> now! <u>Deal</u> of the week! <u>Buy</u> now! Rolex w4tchs - \$100 Medlcine (any kind) - £50 Also low cost M0rgages available.</div>
---	--	---

Building a spam classifier

How to try to reduce your spam classifier's error?

- Collect more data. E.g., "Honeypot" project.
- Develop sophisticated features based on email routing (from email header).
- Define sophisticated features from email body. E.g., should "discounting" and "discount" be treated as the same word.
- Design algorithms to detect misspellings. E.g., w4tches, med1cine, m0rtgage.

2. Error analysis

Error analysis

$m_{cv} =$ ~~500~~⁵⁰⁰⁰ examples in cross validation set.

Algorithm misclassifies ~~100~~¹⁰⁰⁰ of them.

Manually examine 100 examples and categorize them based on common traits.

- Pharma: 21 → more data features
- Deliberate misspellings (w4tches, med1cine): 3
- Unusual email routing: 7
- Steal passwords (phishing): 18 → more data features
- Spam message in embedded image: 5

3. Adding data

Adding data

Add more data of everything. E.g., "Honeypot" project.


Add more data of the types where error analysis has indicated it might help.

Pharma spam
E.g., Go to unlabeled data and find more examples of Pharma related spam.


Beyond getting brand new training examples (x,y) , another technique: Data augmentation

Data augmentation

Augmentation: modifying an existing training example to create a new training example.







Data augmentation by introducing distortions



Data augmentation for speech

Speech recognition example


-  Original audio (voice search: "What is today's weather?")
-  + Noisy background: Crowd
-  + Noisy background: Car
-  + Audio on bad cellphone connection

Data augmentation by introducing distortions

Distortion introduced should be representation of the type of noise/distortions in the test set.



Usually does not help to add purely random/meaningless noise to your data.



x_i = intensity (brightness) of pixel i
 $x_i \leftarrow x_i + \text{random noise}$

Artificial data synthesis for photo OCR



[<http://www.publicdomainpictures.net/view-image.php?image=5745&picture=times-square>]

Artificial data synthesis for photo OCR



Real data



Synthetic data

[Adam Coates and Tao Wang]

Engineering the data used by your system

Conventional
model-centric
approach:

$AI = \text{Code} + \text{Data}$
(algorithm/model)

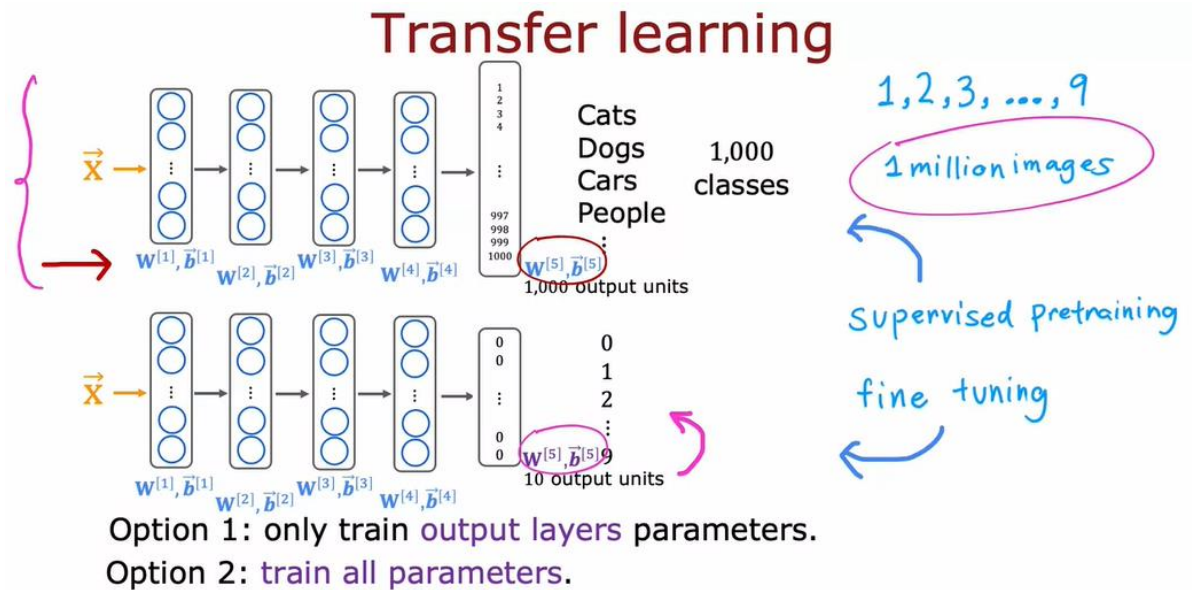
work on this

Data-centric
approach:


$AI = \text{Code} + \text{Data}$
(algorithm/model)

work on this

4. Transfer learning: using data from a different task



Why does transfer learning work?



Transfer learning summary

1. Download neural network parameters pretrained on a large dataset with same input type (e.g., images, audio, text) as your application (or train your own).

1 million images

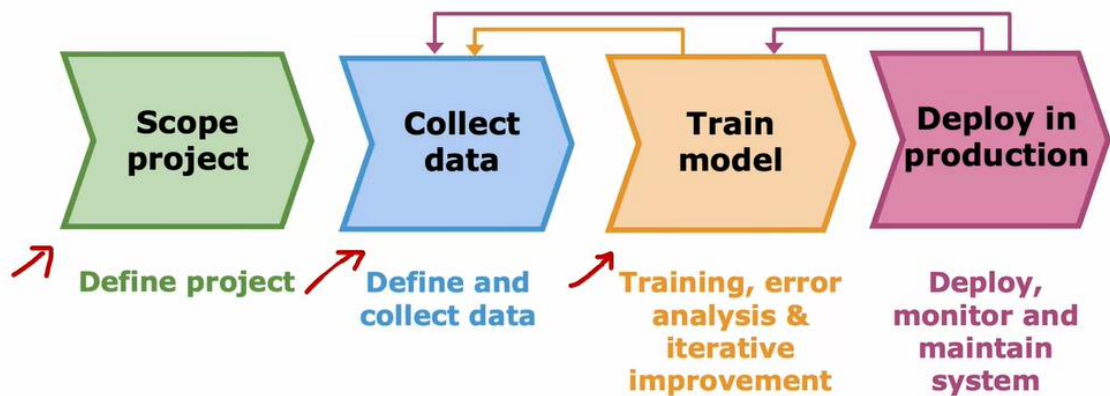
2. Further train (fine tune) the network on your own data.

1000 images

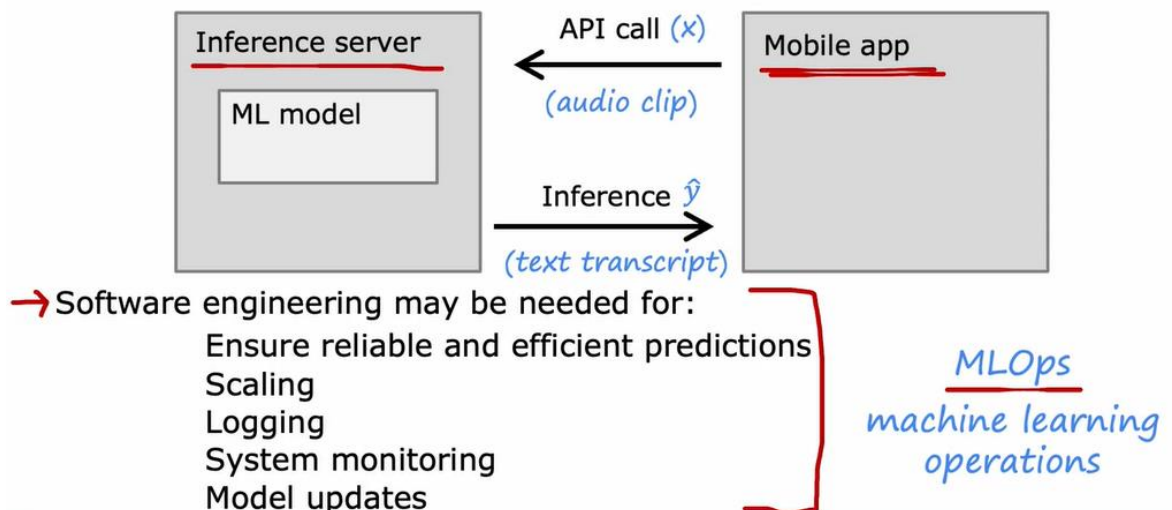
50 images

5. Full cycle of a machine learning project

Full cycle of a machine learning project



Deployment



6. Fairness, bias, and ethics

Bias

Hiring tool that discriminates against women.

Facial recognition system matching dark skinned individuals to criminal mugshots.

Biased bank loan approvals.

Toxic effect of reinforcing negative stereotypes.

Adverse use cases

Deepfakes

Spreading toxic/incendiary speech through optimizing for engagement.

Generating fake content for commercial or political purposes.

Using ML to build harmful products, commit fraud etc.


Spam vs anti-spam : fraud vs anti-fraud.

Guidelines

Get a diverse team to brainstorm things that might go wrong, with emphasis on possible harm to vulnerable groups.

Carry out literature search on standards/guidelines for your industry.

Audit systems against possible harm prior to deployment.



Develop mitigation plan (if applicable), and after deployment, monitor for possible harm.

7. Practice quiz

Which of these is a way to do error analysis?

- ☐ Calculating the test error J_{test}
- ☐ Calculating the training error J_{train}
- ☒ Manually examine a sample of the training examples that the model misclassified in order to identify common traits and trends.
- ☐ Collecting additional training data in order to help the algorithm do better.

✓ **Correct**

Correct. By identifying similar types of errors, you can collect more data that are similar to these misclassified examples in order to train the model to improve on these types of examples.

We sometimes take an existing training example and modify it (for example, by rotating an image slightly) to create a new example with the same label. What is this process called?

- ☐ Machine learning diagnostic
- ☒ Data augmentation
- ☐ Bias/variance analysis
- ☐ Error analysis

✓ **Correct**

Yes! Modifying existing data (such as images, or audio) is called data augmentation.

What are two possible ways to perform transfer learning? Hint: two of the four choices are correct.

- ☒ You can choose to train just the output layers' parameters and leave the other parameters of the model fixed.

✓ **Correct**

Correct. The earlier layers of the model may be reusable as is, because they are identifying low level features that are relevant to your task.

- ☐ Download a pre-trained model and use it for prediction without modifying or re-training it.
- ☒ You can choose to train all parameters of the model, including the output layers, as well as the earlier layers.

✓ **Correct**

Correct. It may help to train all the layers of the model on your own training set. This may take more time compared to if you just trained the parameters of the output layers.

- ☐ Given a dataset, pre-train and then further fine tune a neural network on the same dataset.

Skewed Dataset

1. Error metrics for skewed dataset

Rare disease classification example

Train classifier $f_{\vec{w},b}(\vec{x})$ ($y = 1$ if disease present,
 $y = 0$ otherwise)

Find that you've got 1% error on test set
(99% correct diagnoses)

Only 0.5% of patients have the disease

99.5% accuracy, 0.5% error

`print("y=0")`

1%

1.2%


less

Usefulness

more

Precision/recall

$y = 1$ in presence of rare class we want to detect.



2. Trading off precision and recall

Trading off precision and recall

Logistic regression: $0 < f_{\vec{w},b}(\vec{x}) < 1$
 → Predict 1 if $f_{\vec{w},b}(\vec{x}) \geq 0.5$ (0.7, 0.4, 0.3)
 → Predict 0 if $f_{\vec{w},b}(\vec{x}) < 0.5$ (0.7, 0.4, 0.3)

precision = $\frac{\text{true positives}}{\text{total predicted positive}}$
recall = $\frac{\text{true positives}}{\text{total actual positive}}$


Suppose we want to predict $y = 1$ (rare disease) only if very confident.

higher precision, lower recall

Suppose we want to avoid missing too many case of rare disease (when in doubt predict $y = 1$)

lower precision, higher recall

More generally predict 1 if: $f_{\vec{w},b}(\vec{x}) \geq \text{threshold}$.



F1 score

How to compare precision/recall numbers?

	Precision (P)	Recall (R)	Average	F ₁ score
Algorithm 1	0.5	0.4	0.45	0.444
Algorithm 2	0.7	0.1	0.4	0.175
Algorithm 3	0.02	1.0	0.501	0.0392

`print("y=1")`

~~Average = $\frac{P+R}{2}$~~

F₁ score = $\frac{1}{\frac{1}{2}(\frac{1}{P} + \frac{1}{R})} = 2 \frac{PR}{P+R}$

Harmonic mean